

TEAM 404 · ROUND 2

AutoFleet AI

Autonomous last-mile disruption resolution

Business Model — who pays, how much, and why

Unit economics · pricing · go-to-market

github.com/Himanshi314/AUTOFLEET

Every figure is an assumption unless marked “measured”. The measured ones come from the working system. The market and pricing figures need validating with one real carrier before anyone acts on them.

Every figure here is an assumption unless marked *measured*. The measured ones come from the working system (cost per incident, routing distribution). The market and pricing figures are reasoned estimates that need validation with one real carrier before anyone should act on them. Say that out loud — a judge who knows logistics will respect it more than a confident invented number.

1. What we sell, in one line

Resolved delivery exceptions. Not software seats, not a dashboard — the outcome. A disruption goes in, a completed first-attempt delivery comes out, and no coordinator touched it.

2. Who pays, and out of which budget

We are not sold to an innovation team. We are sold to whoever owns the exception desk's cost line.

Buyer	Title	Why they sign
Primary	Head of Last-Mile / VP Operations at a 3PL carrier	Owns exception-desk headcount, redelivery fuel and SLA penalties
Strong second	COO / Logistics Head at a diagnostic lab chain or pharma distributor	Owns spoiled-sample write-offs — highest cost per exception of any segment
Third	Head of Field Service (appliance, telecom, utility)	Owns technician no-show and revisit costs on ₹500–5,000 jobs

The four budget lines we reduce:

1. Exception-desk headcount (the largest, and the reason it gets bought)
2. Redelivery fuel and courier-hours
3. SLA penalty payouts to enterprise shippers
4. Written-off spoiled stock (cold chain only — but brutal there)

Explicitly not our buyer: quick-commerce (Blinkit / Zepto class). They already automate rider reassignment in-house and would build, not buy. Saying who you *don't* sell to makes the rest credible.

3. Pricing — and why per-seat is structurally wrong

This is the most important decision in the model, so here is the reasoning rather than just the answer.

Per-seat pricing is misaligned with what we do. Standard SaaS charges per user. AutoFleet *removes* users — its whole value is needing fewer coordinators. Per-seat pricing would mean the customer pays us more precisely when we are working worst. The incentives point in opposite directions.

So we charge per resolved exception. We charge for work done. Value and price move together, the buyer's spend scales with their volume rather than their headcount, and the metric we bill on is the same metric we report success on.

The price point

Per exception	
Value delivered (conservative)	₹60–100 — redelivery fuel + courier-hours + partial labour recovery
Value delivered (with full labour + SLA)	₹120–200
Our COGS (measured)	₹0.80 (~\$0.009)
Proposed price	₹8–15, tiered down by volume

At ₹10 we capture roughly **10–17% of the conservative saving** and leave the rest with the customer. That is an easy conversation: "we keep a tenth, you keep nine."

Structure — three components:

Component	Amount	Why it exists
Platform fee	₹1.5–4L/month by tier	Integration, support, uptime commitment. Makes revenue predictable and stops tiny accounts being unprofitable.
Per resolved exception	₹8–15, volume-tiered	The real meter. Only charged when the system resolves autonomously.
Escalations: free	₹0	If we hand it to a human, we didn't do the job. Don't charge for it.

That last line is the one to say in the pitch. It is commercially cheap (escalations should be a small minority) and it makes the incentive unmistakable: we only get paid when we actually replace the coordinator.

4. Unit economics

Per resolved exception, at ₹10:

	₹	Note
Revenue	10.00	
AI inference (measured)	(0.80)	~\$0.009 blended, after routing
Infra, comms (SMS), logging	(0.60)	estimate
Gross profit	8.60	
Gross margin	~86%	

Account-level, per month

Customer size	Exceptions/day	Platform fee	Usage	MRR	ARR
Large national carrier	20,000	₹4L	₹60L	₹64L	₹7.7Cr (~\$900k)
Mid regional carrier	3,000	₹2L	₹9L	₹11L	₹1.3Cr (~\$155k)

Customer size	Exceptions/day	Platform fee	Usage	MRR	ARR
Diagnostic lab chain	500 (cold chain)	₹1.5L	₹2.3L ¹	₹3.8L	₹46L (~\$54k)

¹ Cold chain prices at a premium (₹15) — a spoiled sample costs far more than a redelivered parcel, so the value per exception is higher.

Why the margin *improves* over time — the unusual bit

Most AI products have COGS that scale linearly with usage, which crushes gross margin as you grow. Ours doesn't, because of the **severity router**.

Measured: the router resolves **11% of incidents with zero AI calls** and another **61% with four instead of six** — cutting blended inference cost from ~\$0.068 to ~\$0.009 per incident, about **7×**.

Every improvement to the router is a permanent margin gain. And as the resolution- policy library grows, more cases become mechanical, so more incidents route to zero AI calls. **The product gets cheaper to run the better it gets.** That is a genuinely attractive property and most AI startups cannot claim it.

5. Revenue evolution — three phases

Nobody buys autonomous customer communication on day one. The commercial model has to mirror the trust ladder.

Phase	What we deliver	What we charge	What we're really buying
1 • Shadow (4–8 weeks)	We watch real exceptions, decide, act on nothing . Log our decision beside the coordinator's.	Free , or a small fixed pilot fee	The only evidence that matters: " <i>we matched or beat your coordinator on N real exceptions.</i> " Also: their exception taxonomy, which becomes our policy library.
2 • Assisted → Guarded (3–6 months)	We propose; the coordinator approves. Then we act alone below a value threshold.	Per-resolution, discounted (₹5–8)	Approval rate as the proof metric. Every override becomes a policy-library entry.
3 • Autonomous	Full resolution within agreed guardrails	Full price (₹8–15) + platform fee	Expansion: more exception types, more cities, more of their network

Shadow mode is the entire sales motion, and it is nearly free to deliver. It needs read-only access — no approval to *act*, which is the hard approval to get. That keeps customer acquisition cost low and gives the sales cycle a natural, data-backed proof point instead of a demo.

6. Market size — bottom-up, not top-down

We are deliberately not quoting a "\$50B logistics market" figure. Bottom-up by account is more defensible and more useful.

India, addressable accounts:

Segment	Accounts	Realistic ACV	Segment potential
National 3PL / express carriers	~15	₹1.5–7.7Cr	~₹30–60Cr
Regional carriers	~80	₹40L–1.3Cr	~₹50–80Cr
Diagnostic lab / pharma cold chain	~40	₹30L–1Cr	~₹15–25Cr
Enterprise field service	~150	₹20–60L	~₹40–60Cr

Serviceable India opportunity: roughly ₹135–225Cr ARR (~\$16–27M). Beachhead first: **five mid-size carriers = ~₹6Cr ARR (~\$700k)**, which is a real company.

Same playbook extends to South-East Asia, the Middle East and LATAM, where last-mile density and manual coordination look similar. Do not put that in the first pitch — beachhead credibility beats a big number.

7. Costs to run

Monthly, early stage	
AI inference	Scales with usage — ~₹0.80/exception (<i>measured</i>)
Cloud infra	₹40–80k (the app itself is light: stdlib Python, no DB, no build step)
Integration engineering	1–2 engineers — the real cost centre
Sales / customer success	1 person until ~5 accounts

The honest cost warning: our cost is not compute, it's **integration**. Every TMS is different. Each new carrier needs connectors to their dispatch system, driver app and comms provider. That is what will consume the team, and it is also what compounds into a moat — but budget for it as engineering time, not as licence cost.

8. Moat — honest version

What is not a moat: the agent chain. An incumbent could rebuild it in a weekend. Do not claim it as defensibility; a technical judge will call it.

What actually compounds:

1. **The resolution-policy library.** Every shadow-mode disagreement teaches the system something a competitor starting today doesn't know. This grows with usage and cannot be copied from the outside.
2. **Integration surface.** Connectors into a dozen TMS variants are unglamorous, slow, and a genuine barrier.
3. **The trust record.** "*We matched human coordinators on 40,000 real exceptions*" is the asset no new entrant has, and it is what closes the second customer.
4. **Switching cost.** Once we are in the exception path, removing us means rebuilding the exception desk.

9. Risks, and what would kill this

Naming these makes the model credible rather than weaker.

Risk	Severity	Response
Carriers build it in-house. The big ones have engineering teams.	High	Sell to mid-size carriers who won't build. Be the layer, not the platform.
Rules handle most of it — <i>measured</i> : 66% of exceptions have one obviously best driver, so a rules engine covers the head of the distribution	High	We already route those to zero AI calls. Our price reflects resolution, not intelligence — we're paid for the outcome regardless of method. But the value case rests on the tail being large and expensive, and we have not proven that with real data.
Trust / adoption is slow. Ops directors don't hand over customer comms.	Medium	The three-phase ladder is designed for exactly this. Long sales cycle is a fact to plan for, not a bug.
One bad autonomous decision damages a carrier's key account	Medium	Hard guardrails already built: no payments, no cancellations, value thresholds, and honest escalation.
Inference price rises	Low	The router already cut cost 7×; margin has room.

10. The 30-second version

We sell resolved delivery exceptions, not software seats — because our whole value is that you need fewer coordinators, so charging per seat would point the incentives the wrong way. Roughly ₹10 per exception we resolve autonomously, free when we have to escalate to a human. It costs us about eighty paise to resolve one, so gross margin is around 86% — and it *improves* as our router learns to resolve more cases without waking a language model at all. We land through free shadow mode, where we decide but don't act, because the only thing that sells this is proof we match a real coordinator on real exceptions. Five mid-size carriers is about ₹6 crore of ARR.

Appendix — where each number comes from

Figure	Source
₹0.80 (~\$0.009) COGS per exception	Measured. 6,120 input + ~1,500 output tokens at Opus 5 rates, blended across the routing distribution
7× cost reduction from routing	Measured. \$0.068 → \$0.009 blended
29% full / 61% partial / 11% deterministic	Measured across all 28 disruption × delivery combinations
66% of exceptions have a decisive winner	Measured on the demo fleet (ranker margin ≥ 0.10). <i>Not</i> a production distribution
14 coordinator-minutes per incident	Assumption — midpoint of the stated 10–20 min manual window
₹60–200 value per exception, and ₹30–60/parcel carrier margin	Assumption, from general commentary on Indian express economics. Verify both against one real carrier's cost data before quoting either.
₹8–15 price per exception	Assumption. Derived to capture 10–17% of the conservative value estimate
Account counts by segment	Assumption. Order-of-magnitude reasoning, not a researched census