

Run of Show

Six beats, just under five minutes. Every click here was rehearsed against the live build, so the numbers you point at are the numbers that actually show up. Read the pre-flight first — two of those steps are the difference between a demo and a retake. The two **business** lines are marked; they are the only places you talk about money or regulation.

4:50 runtime 6 beats 2 Resets mid-take 3/min incident ceiling **Numbers stay in English** — always

PRE-FLIGHT — SIX THINGS BEFORE YOU HIT RECORD

- 1 **Check the model is actually alive.** Run it, read the last line.

```
python -m autofleet.llm

model      openai/gpt-oss-120b
probing    ... OK
VERIFIED - agents will use the model.
```

If it says NOT VERIFIED, stop right there. The demo will look completely normal while every agent runs on deterministic text — this is the one failure that is invisible on screen.

- 2 **Start the server,** open the dashboard.

```
python -X utf8 server.py --port 8600    →    http://127.0.0.1:8600
```

Corner must read **LIVE · openai/gpt-oss-120b · verified**. Bas, that's the only check.

- 3 **Hit Reset,** then confirm the fleet clock reads **17:2x** and every impact tile is zero.

The sim clock crosses midnight over a long session. Telling an evening-peak story with a clock showing 07:57 is exactly the detail a judge picks on.

- 4 **Browser at 100% zoom, window roughly 1500×950.**

Below ~780px of height the live row gets squeezed and the map turns into a strip. Full screen on your laptop is fine.

- 5 **Notifications off, every other tab closed.**

One Teams toast over the agent feed and the take is gone.

6 One dry run of Beat 4 — the conflict check — then Reset again.

The whole submission rests on that beat, and a dry run warms the model so your first real take isn't the slowest one.

0:00

25 sec

A bike just died. Ab kya?

Operations view, nothing triggered. Let the board move behind you.

LAND THIS LINE

**"AutoFleet does it without the phone call —
and every decision it takes, you can audit."**

DO

Nothing. Cursor still. Let the ETAs tick and the clock move.

SAY

So here's the scene. A courier's bike dies eight kilometres from the drop. What happens today? Someone in a control room notices, starts calling around, re-plans it by hand. Ten to twenty minutes — and this happens thousands of times a day.

AutoFleet does it without the phone call. And every decision it takes, you can audit. Bas.

JUDGE SEES

- A live board — ETAs changing, fleet clock moving. Not a mockup.
- **LIVE · openai/gpt-oss-120b · verified** in the corner.

0:25

70 sec

The rider speaks. Six agents listen.

Longest beat, and it earns everything after it. Don't rush the Input block.

LAND THIS LINE

"This isn't one model writing six paragraphs. Dekho — this is what the Customer Agent actually received."

DO

1. **Courier** tab → pick **Priya · DR-01**.
2. Scroll to **Report a problem** → tap  **Vehicle Breakdown**.
3. Straight back to **Operations**.
4. When the chain finishes, expand **Input** on the **Customer Agent** card.
5. Point at the green **fact-check** badge on the Coordinator card.

SAY

The rider raised this, not the system — because he knows the bike is dead before any sensor does. Corridor gridlock is the opposite: that he can't see from the saddle, toh operations raises it instead.

Six agents woke up. Nobody typed a prompt.

And this isn't one model writing six paragraphs. Dekho — this is what the Customer Agent actually received. The Risk Agent's words, exactly as they were. Zero, one, two, three prior decisions as you go down the chain. Matlab they're genuinely passing work to each other.

Then the last agent's summary gets fact-checked against the numbers the models computed. Four claims, all verified. If it had invented a number, this badge would say so.

JUDGE SEES

- `trigger=courier` in the log — a field report, not a console click.
- Six agent cards streaming, each with a **received N** count that grows.
- A real tool call, reassigning to a named person.
- **Human interventions: 0** — and that number is counted, not hardcoded.

If a card says "simulated" — model is rate-limited or the key failed. Kill the take, wait 60 seconds, re-run `python -m autofleet.llm`. Never narrate over fallback text.

1:35

30 sec

Knowing when *not* to use AI

Short beat, but this is where engineering maturity shows.

LAND THIS LINE

"Everyone's pitch is 'we used AI'. Ours is that we measured where it wasn't needed — and took it out."

DO

1. **Reset.** Wait two seconds.
2. On **D-102**, hit  **Recipient Not Home.**
3. Point at the router line in the log.

SAY

Different problem — and now only four of six agents run. Nobody's being reassigned, toh the Resource and Delivery agents have nothing to decide. And a role with nothing to decide doesn't get woken at all.

Everyone's pitch is "we used AI". Ours is that we measured where it wasn't needed and took it out. Seedha. Two model calls saved on this one incident — and it's a deterministic branch, not a judgement call.

JUDGE SEES

- `routed as 'partial chain' · 4/6 agents will run · 2 model call(s) avoided`
- Two agent cards greyed out, with the reason they were skipped.

2:05

90 sec

The system says no — and asks a human

The round-2 capability, and the centre of everything. **D-104 escalates every single time after a Reset** — rehearsed, not luck.

LAND THIS LINE

“Every option breaks somebody's stated goal — toh it doesn't commit. It stops, and it asks.”

DO

1. **Reset.** Click **Show all** on the intent register.
2. On **D-104 · Vikram Shetty · Whitefield**, hit  **Vehicle Breakdown.**
3. Let the **Intent conflict check** card render. Pause on it.
4. Scroll to the amber **needs a decision** panel.
5. Click **Override Operations's intent.**
6. Point at the **Decisions taken** line underneath — **stay on it** while you say the regulation line.

SAY

These are goals people actually stated. Not constraints we assumed — matlab things a named person said, in their own words, and can withdraw. Ananya wants nothing after six. Priya wants thirty minutes at the end of her shift. Operations capped empty running at eight kilometres.

Now watch it refuse. Both available couriers are more than seventeen kilometres from the payload, against an eight kilometre limit that Operations set. Every option breaks somebody's stated goal — toh it doesn't commit. It stops, and it asks.

And here's what a human gets handed: what it was about to do, whose goal breaks, the arithmetic behind that — 17.7 against 8, so 9.7 over — and the options that are genuinely open.

I'll lift the ceiling for this one case. The system re-decides with the new fact. It wasn't overruled — it was re-run. And the override is on record, with my name on it.

BUSINESS LINE.

And this isn't just good practice. In the EU, AI that allocates work to people is regulated as high-risk — it needs logging, human oversight, and a record of who decided what. That panel you just watched *is* that record.

JUDGE SEES

- 3 stated intents checked against 2 options · 0 clear, 2 blocked
- Every option listed with **what passed as well as what failed** — that's how you know the check actually ran.
- Evidence chips: approach km 17.68 · limit km 8.0 · margin km -9.68
- **Human interventions: 1** — the autonomy number stays honest.
- An audit line naming who overrode whose goal, and when.

If D-104 resolves instead of escalating — you triggered before the Reset landed, or the board had already respawned. Reset, wait three seconds, retry on D-104. **Backup:** any delivery carrying a **tricky** chip escalates about 40% of the time — use one of those.

If a judge asks which regulation — do not quote an article number. Say: *"I'd want to confirm the exact article before quoting it, but the obligation is logging, human oversight and traceability for task-allocation systems."* That is stronger than a wrong citation, and someone in that room may know this law better than you do.

3:35

30 sec

And the banda whose bike died?

This is the beat they remember eight hours later. Say it simply, don't oversell.

LAND THIS LINE

"Everyone optimises for the customer or the company. The rider? Nobody instruments the rider."

DO

1. **Courier** tab → pick the rider whose bike broke (red dot).
2. Then pick whoever took the job over; scroll to **Your route**.

SAY

The rider whose bike died: roadside assistance dispatched, earnings for the leg he completed protected, no reliability penalty logged. Galti uski nahi thi.

Everyone optimises for the customer or the company. The rider? Nobody instruments the rider. And whoever picked the job up gets two legs — ride to the payload, then carry it — drawn from the same coordinates the ETA came from.

JUDGE SEES

- Support dispatched · **earnings protected** · no penalty.
- A two-leg route: dashed to the payload, solid to the recipient.

4:05

45 sec

Proof, the business, phir seedhi baat

Three things, in this order: the toggle proves it isn't cosmetic, then who pays, then the limits. Don't reorder — the limits have to land last.

LAND THIS LINE

**"That's the difference between a capability and a caption
— and you can do it yourself, on our machine, right now."**

DO

1. **Show all** → switch off **Operations' 8 km ceiling**.
2. Reset, then the same  **Vehicle Breakdown** on **D-104**.
3. Let it resolve. Stop recording on the impact strip.

SAY

Same incident. One goal withdrawn. Now it doesn't stop — it commits a courier. That's the difference between a capability and a caption, and you can do it yourself, on our machine, right now.

BUSINESS LINE.

Quick word on the business. We don't sell software seats — our whole value is needing fewer coordinators, so per-seat pricing would mean earning more when we work worse. We sell resolved exceptions. Carriers already spend this money, just on coordinator salaries and redelivery fuel. Five mid-size carriers is roughly seven hundred thousand dollars a year.

And seedhi baat: the geography is real Bengaluru coordinates, haversine distance, a documented circuitry factor. The telemetry is simulated, because we don't have a fleet. Every figure on that strip is labelled — real, sourced estimate, or assumption — on the Assumptions page. 116 tests. Thank you.

JUDGE SEES

- The same trigger producing a **different outcome**, because one goal changed.
- A pricing model aligned with the value — not a per-seat number pulled from air.
- A team naming its own limits before anyone asks.

DO NOT, DURING THE TAKE

Don't open /react. It has never been verified running a full incident. The vanilla dashboard is the product.

Don't spam the trigger buttons. Roughly three incidents a minute is the token ceiling. Past that, agents drop to deterministic text mid-chain and the badge changes while you're talking.

Don't arm Autonomous and wait for it. It fires on real risk every six or seven minutes. Arm it at the start and let it land during questions — “nobody clicked that” beats dead air.

Don't leave it running 20+ minutes without a Reset. The clock crosses midnight and the evening-peak story stops making sense.

LIMITS — SAY THESE BEFORE YOU'RE ASKED

It's a simulation

Real coordinates, real distance maths, simulated telemetry. Say it first and it becomes a credibility signal instead of a hole.

The model is hand-calibrated

Seven features, stated weights, fully decomposable. No labelled failure data — faking a training run would be worse.

Whole-delivery handover

We don't model multi-drop, so one courier carries one live job. The load tile is context, not a live constraint.

No shift-start awareness

A courier free in ten minutes is invisible to the ranker today. Known gap, not a surprise.

No integration, no pilot

No carrier integration, no rider interviewed yet, no trained model. Naming the gaps stops a judge enjoying finding them.

Revenue today is zero

No pilot, no integration, no rider interviewed. The business numbers are a model, not a forecast. Say it flatly and move on.

Escalation is rare by design

Usually the best courier is also the acceptable one. The gate exists for when they diverge — and we can show you both.

Rehearsed on the live build: Groq `openai/gpt-oss-120b`, 6/6 agents on the model, D-104 escalating on the approach ceiling at 17.68 km against 8.0 km. `python run_tests.py` → 116 passed.